

Towards Testing ACID Compliance in the LDBC Social Network Benchmark

Jack Waudby, Benjamin A. Steer, Karim Karimov, József Marton,
Peter Boncz, and Gábor Szárnyas



Hungarian Academy of
Sciences



Linked Data Benchmark Council



LDBC is a non-profit organization established in 2012, dedicated to defining benchmarks and *auditing results* for graph data management software.

Similar goals to TPC's in the relational domain:

- facilitate fair comparison
- drive competition
- capture an understanding of the field



LDBC Structure



Task Forces;

- Social Network Benchmark
- Semantic Publishing Benchmark
- Graphalytics

Working Groups;

- Property Graph Schema
- Query Language (G-CORE)
- Formal Semantics (GQL)
- Existing Languages (literature review)

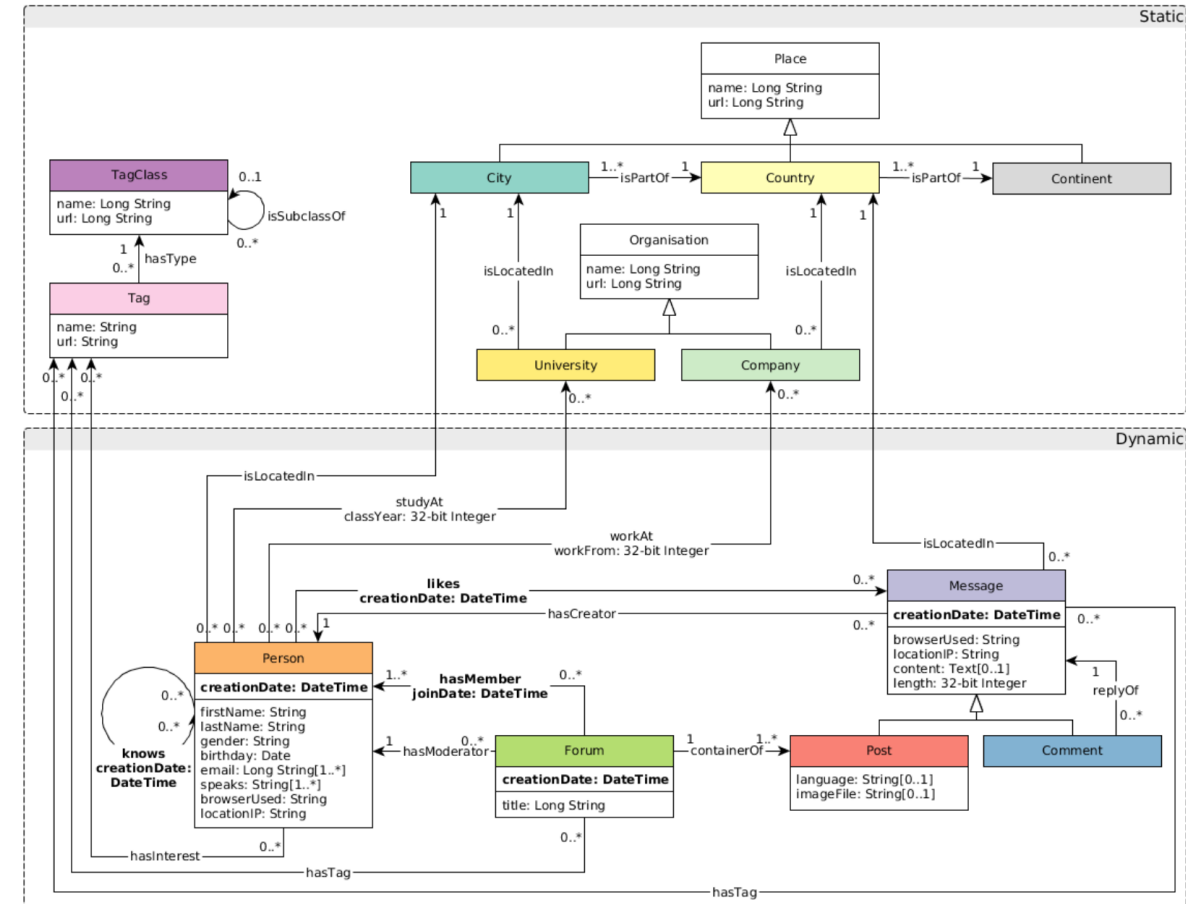
Membership;

- Non-profit orgs (ENS-Paris, University of Edinburgh)
- Commercial orgs (TigerGraph, Neo4j, Oracle, OpenLink Software, Ontotext, JCC Consulting, Intel)
- Individual members
- Associate members



LDBC Social Network Benchmark (SNB) Suite

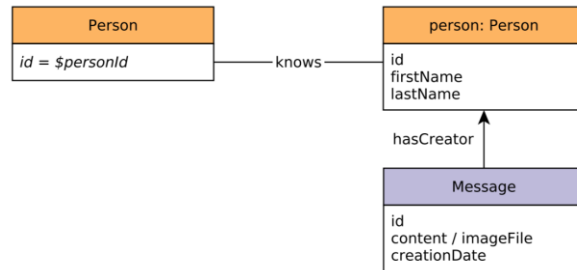
- Models a social network, e.g., Facebook
- People connect with each other and post messages in groups
- Correlated synthetic data produced by LDBC's Datagen
- 2 workloads; Interactive and Business Intelligence (BI)



LDBC Social Network Benchmark (SNB) Suite

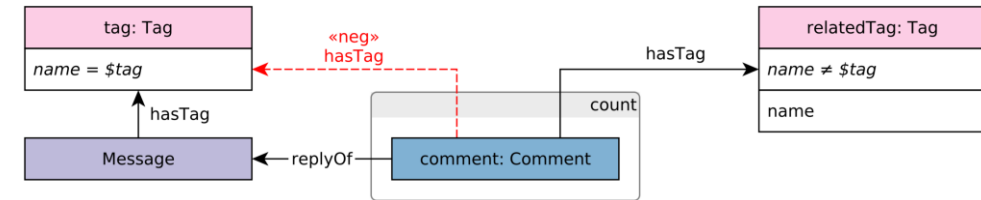
- **Interactive (OLTP)**

- complex/short read queries, updates
- explore neighborhood of a node/query
- established, multiple implementations



- **BI (OLAP)**

- currently read-only queries
- large portion of the graph/query
- ongoing adoption



Erling et al.
The LDBC Social Network Benchmark: Interactive Workload,
SIGMOD 2015



Szárnyas et al.
An early look at the LDBC Social Network Benchmark's Business Intelligence Workload,
GRADES-NDA 2018



Motivation

- Starting receiving requests for audits;
 - first audit completed July 2020, FMA's TuGraph
- ACID compliance important for fair comparison between systems
- No mechanism for validating ACID compliance*



- Design an ACID compliance test suite
- Focus on Atomicity and Isolation

**Durability test already part of benchmark specification*



Related Work

- **TPC-* Tests:**
assume lock-based concurrency control, tests for 3 isolation anomalies
→ not generalizable, limited anomaly test coverage
- **Hermitage (Martin Kleppmann)**
tests performed by hand
→ hard to induce anomalies that relied on fast timings
- **Jepsen (Kyle Kingsbury)**
focuses on distributed systems under various failure modes
→ too heavyweight



Design Considerations

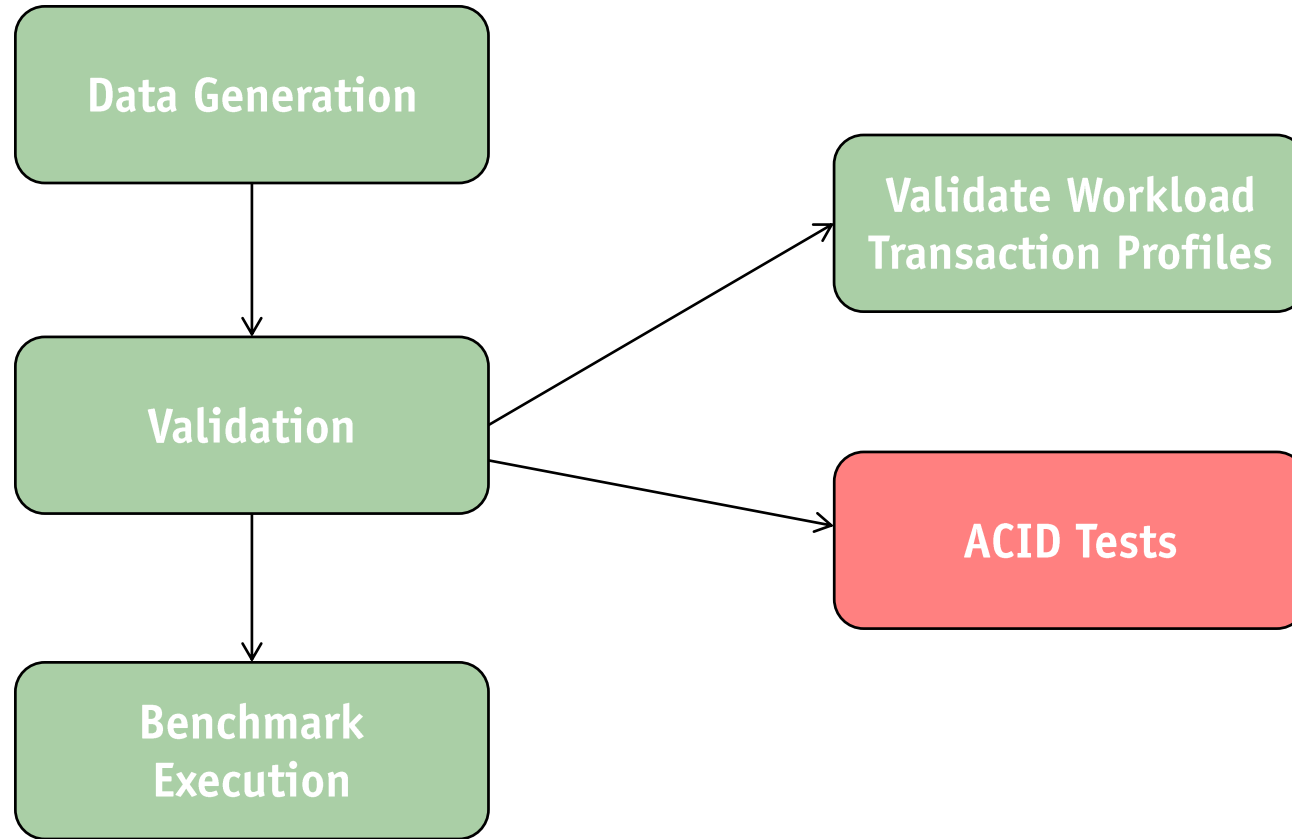
Disclaimer: verifying ACID compliance with a finite set of tests is not possible

1. **Generalizable** – agnostic of system-level implementation details and query API
2. **Lightweight** – not add significant overhead to benchmarking process
3. **Improved Coverage** – test for more isolation anomalies, e.g., lost updates, write skew



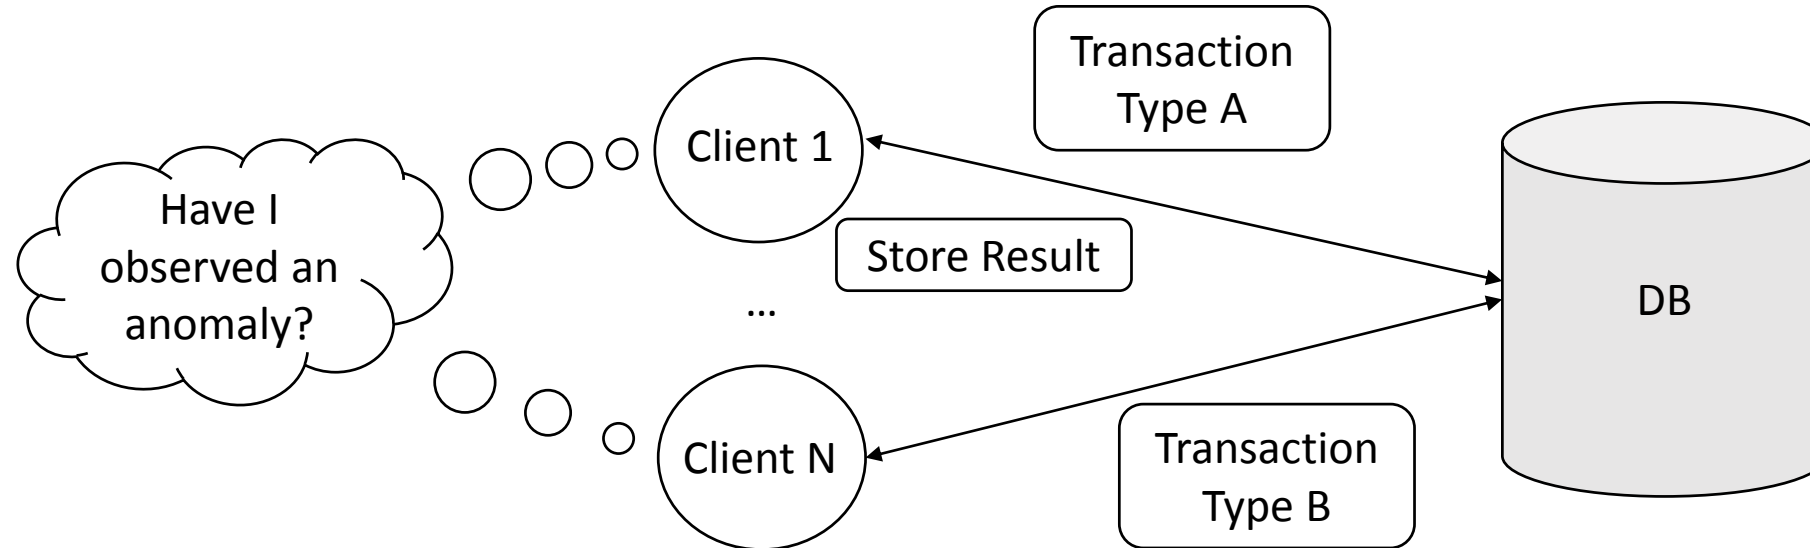
LDBC ACID Test Suite

Included in benchmark workflow as an additional step in the **Validation Phase**



LDBC ACID Test Suite Design

- Based solely on *client observations* to detect anomalies (generalizable)
- Each test consists of a handcraft set of transactions, which when interleaved create conditions in an anomaly could occur
- After execution, transaction results are gathered, and an *anomaly check* performed



LDBC ACID Test Suite Execution Flow

For each test;

1. Load required test graph
2. Initiate N clients
3. Execute test's transaction set for duration T
4. Gather transaction results from clients
5. Perform anomaly check



Test Coverage

- Atomicity;
 - Commit
 - Rollback
- Isolation;
 - Dirty Writes¹, Aborted Reads¹, Intermediate Reads¹, Circular Information Flow¹, Write Skew¹, Lost Updates¹, Item-Many Preceders², Predicate-Many Preceders², Observed Transaction Vanishes², and Fractured Reads³.



Adya et al.¹
Generalized Isolation Level Definitions, ICDE 2000



Bailis et al.²
Highly Available Transactions: Virtues & Limitations, VLDB 2014




Bailis et al.³
Scalable Atomic Visibility with RAMP Transactions, SIGMOD 2014



Example: Lost Updates

A **lost update** occurs when 2 transactions concurrently attempt to make conditional modifications to the same data item(s)



If T1 and T2 are executed sequential, $X=7$, T2's update is "lost"



Example: Lost Updates

1. Load a test graph containing *Person* nodes, with unique *id* and *numFriends* properties
2. Clients choose a random *Person* and increment its *numFriends* property
3. Clients store local counters (*expNumFriends*) for each *Person*, which is incremented each time a *Person* is selected and the transaction successfully commits
4. After execution, gather counters, then,
5. Perform anomaly check: for each *Person*,
(global) *expNumFriends* = (observed) *numFriends*

Person
<i>\$id</i>
<i>\$numFriends</i>

```
MATCH (p:Person {id: $personId})  
SET p.numFriends = p.numFriends + 1
```



Experimental Setup

- Implemented as extensible framework in a Java application
- Ubuntu 18.04 running AdoptOpenJDK 11.0.4.hs
- 4 graph database systems and 1 relational database;
 - Neo4j 3.5.20 and 4.1.1
 - Memgraph 1.0
 - Dgraph 20.07.0
 - JanusGraph 0.5.2 (BerkeleyDB 7.5.11 and Cassandra 3.11.0 backends)
 - PostgreSQL 9.6
- For all systems, we used their declarative query languages (Cypher, GraphQL, Gremlin, and SQL) and the officially recommended Java drivers



Results

- Overall, most systems met their claims
- Selected results;
 - Memgraph's default isolation level is Snapshot Isolation; only Write Skew occurred
 - Neo4j's default isolation level is Read Committed with *some* built-in protection against Lost Updates; v3.5.20 and v4.1.1 met requirements for Read Committed, but v4.1.1 displayed Lost Updates
 - JanusGraph had the worst user experience, passing some tests due to serving *extremely* stale reads; also, execution of some tests timed out
 - Dgraph's default isolation level is Snapshot Isolation; passed our Write Skew test



Future Work

- Included test suite in LDBC Auditing Policies;
 - Used test suite in recent audit of FMA's TuGraph
- Extend test suite to incorporate complex consistency constraints;
 - Graph databases generally do not support constraints; sometimes domain and primary key constraints
 - Graph-specific constraints are expected to be introduced;
 - (Partial) compliance to a schema on top of property graphs,
 - Structural constraints, e.g., connectedness of the graph, absence of cycles, or arbitrary well-formedness constraints
 - LDBC PGSWG actively working in this area
- Add tests for distributed graph processing systems

